

case_study_part_1_Rondell-Copy2

February 16, 2019

1 Problem 1 from MS Case Study

1.0.1 Rondell King

Background You work at a magic potato trading company. The company earns fees by storing potatoes for clients as well as finding buyers and sellers of potatoes on behalf of clients. The storage fees earned by the company are based off the total market value of the potatoes stored by the clients. These potatoes are magic- they do not spoil while in storage, they can be shrunk so they take up very little space, and the delivery of the potatoes does not incur any cost. Potatoes must be kept separate from each other and therefore the cost to store different potato types may differ. Storage rates will be determined by traders and based off of numerous factors.

You work on the data science team and your role is to work with data to create reports, models, and make the business smarter and more efficient Two parts Analyse the given sets of data and answer the questions below. Help the management team understand what is happening in the potato market and what client activity looks like. Create a web tool that helps capture and save potato prices seen in the market

Part 1 Background You are given the Data_Files excel sheet which contains information about potatoes, the company's clients, and a snapshot of client potato positions for a period of time. A quantity of Null represents no position. Management has some questions regarding these data and would like to know your interpretations. You will present your findings at the quarterly management meeting.

Questions * Which clients have the largest potato stockpile based on market value? * Which clients are the most active? * Which potatoes are most activity traded? * What client activity trends do you see? * What potato price trends do you see? * Are there any factors that can help predict potato prices? * Are there any factors that can help predict client activity? * Summarize what has happened during this period

2 First Step is to load the excel sheets using Pandas.

Examine the data and check for missing values in the tables.

```
In [25]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```

###
### LOAD AND SANITIZE DATA###
###

# Load xls sheets into dataframes
data_file = '2_Data_Files.xlsx' # File must be in the same directory as this script

# Each sheet will be stored in the dict as a dataframe object
df_dict = pd.read_excel(data_file, sheetname=['Potatoes_Info', 'Client_Info', 'Potatoes
positions = df_dict.get('Potatoes_Positions') # Trade data from 5/1/2020 to 8/31/2020.
client_info = df_dict.get('Client_Info') # Key = Client ID
potato_info = df_dict.get('Potatoes_Info') # Ket = Product ID

# Merge the sheets into one dataframe object.
positions_merged = positions.merge(client_info, on='Client ID')
positions_merged = positions_merged.merge(potato_info, on='Product ID')

# Sanitize the Data by replacing [NULL] with 0 (A value of [NULL] represents no position)

print(positions_merged.isnull().sum())#213 prices are missing.

```

/anaconda3/lib/python3.7/site-packages/pandas/io/excel.py:329: FutureWarning: The `sheetname` ke
 **kwds)

```

Date                                0
Client ID                           0
Product ID                           0
Price                               213
Quantity                             0
First Name                           0
Last Name                            0
Client Type                           0
Client Location                       0
Variety Name                          0
Country                              0
Shape of tuber                        0
Colour of skin                        0
Colour of flesh                       0
Depth of eyes                         0
Smoothness of skin                   0
Colour of base of lightsprout         0
Maturity                             0
Height of plants                      0
Frequency of berries                  0
dtype: int64

```

2.0.1 Convert Null prices to 0

```
In [26]: positions_merged['Quantity'].replace(' [NULL]', 0, inplace=True)
```

3 Which clients have the largest potato stockpile based on market value?

Top 3 Clients in terms of market value are:

1. Destiny
2. Jamel
3. Samuel

I found the top market value by adding a new column to the dataframe.

Price X Quantity = Market Value

Sorting the table by Market value shows the clients with the highest Totals.

```
In [27]: # Add new column for Market value => Price * Quantity.
positions_merged['Market Value'] = positions_merged.Price * positions_merged.Quantity
positions_merged['Market Value'].fillna(0,inplace=True)

# Aggregate on client ID to find the clients with highest market value.
top_market_value = positions_merged.groupby(['Client ID',
                                             'First Name ',
                                             'Last Name'], as_index=False)['Market Value'].max()
top_market_value = top_market_value[['Client ID', 'First Name ', 'Last Name', 'Market Value']]
top_market_value.sort_values('Market Value', ascending=False)

print(top_market_value.to_string(formatters={'Market Value': '${:,.2f}'.format}))
```

	Client ID	First Name	Last Name	Market Value
4	27	Destiny	Aldridge	\$227,333,258,486.22
7	38	Jamel	Caruso	\$4,099,865,425.30
21	82	Samuel	Cheney	\$381,338,437.50
18	74	Nakisha	Southerland	\$253,565,363.06
5	30	Eulalia	Culver	\$178,615,864.61
22	89	Tammera	Lassiter	\$77,000,752.74
9	44	Jeremy	Kong	\$68,287,532.35
13	50	Landon	Kyle	\$47,785,223.01
23	94	Valeri	Burchfield	\$47,421,539.53
0	17	Chan	Vue	\$38,883,442.46
12	49	Kirstin	Browne	\$34,918,354.53
14	52	Larita	Albright	\$27,403,986.45
16	64	Mariko	Giles	\$20,889,961.86
19	77	Rasheeda	Spivey	\$16,789,133.46
20	78	Rashida	Proctor	\$6,808,320.95
2	21	Chase	Woody	\$1,765,570.00

3	24	Daine	Gustafson	\$1,068,943.80
15	59	Loraine	Mcdermott	\$995,961.26
6	33	Georgiana	Moya	\$432,154.95
17	65	Maura	Jeffrey	\$359,841.62
8	41	Jeane	Peoples	\$336,669.20
10	45	Kazuko	Steward	\$105,272.75
11	47	Kiera	Staton	\$0.00
1	18	Chantelle	Hollins	\$0.00

4 Which clients are the most active?

Count how many times the client name is in the positions data. Get a count of unique values for First Name

Below are the are names of the 10 most active clients and their total trades during this period.

```
In [28]: client_acitivy = positions_merged.groupby(['Client ID',
                                                    'First Name ',
                                                    'Last Name'])['Date'].count()
client_acitivy.sort_values(ascending=False).head(10)
```

```
Out[28]: Client ID  First Name  Last Name
38              Jamel      Caruso      10239
74             Nakisha  Southerland    3575
50             Landon      Kyle       2091
94             Valeri   Burchfield    1729
64             Mariko      Giles     1636
44             Jeremy      Kong     1629
27             Destiny  Aldridge    1202
17              Chan      Vue       961
89             Tammera  Lassiter     879
30             Eulalia   Culver      740
Name: Date, dtype: int64
```

5 Which potatoes are most actively traded?

30 Most actively traded potatoes and their trade count.

```
In [29]: potato_trade_activity = positions_merged.groupby(['Product ID',
                                                            'Variety Name'])['Date'].count()
potato_trade_activity.sort_values(ascending=False).head(30)
```

```
Out[29]: Product ID  Variety Name
16              Argos      724
172             Lorimer     682
45             British Queen  615
93             Edzell Blue   611
```

10	Amorosa	582
19	Arran Banner	548
281	Shepody	514
100	Emma	511
14	Anya	482
7	Ambassador	480
251	Rocket	474
2	Accord	471
28	Atlantic	461
57	Carnaval	422
102	Erntestolz	404
73	Colleen	401
98	Emblem	384
218	Pentland Dell	383
15	Apache	382
145	Jubilee	377
122	Gourmandine	372
36	Barna	358
217	Pentland Crown	354
278	Shannon	343
81	Cultra	339
17	Ariata	333
254	Roscor	319
23	Arrow	317
83	Desiree	308
230	Pizazz	305

Name: Date, dtype: int64

6 What client activity trends do you see?

Client activity is the lowest during May and peaks in June. Throughout August begins to taper out. Seasonality plays a role here, where it appears the market for potatoes experiences peak activity during the summer season.

In August client activity is relatively low when compared to Jun and July but there is much less volatility in activity during August. Between May and July, trade volume swings are much more intense.

Highest trade activity count:404 occurs on 2020-06-12

Lowest trade activity count:232 occurs on 2020-05-05

```
In [41]: # Max and min trade activity
import seaborn as sns
sns.set(rc={'figure.figsize':(30,20)})
max_trds = positions_merged['Date'].value_counts().max()
min_trds = positions_merged['Date'].value_counts().min()
max_date = positions_merged['Date'].value_counts().idxmax()
min_date = positions_merged['Date'].value_counts().idxmin()
```

```

print("Highest trade activity count:{} occurs on {}".format(max_trds, max_date))
print("Lowest trade activity count:{} occurs on {}".format(min_trds, min_date))

activity_plt = positions_merged['Date'].value_counts().plot(title = 'Trade Activity')
activity_plt.set_xlabel('Trade Date', fontsize=15)

client_activity_table = positions_merged.set_index('Date')

```

Highest trade activity count:404 occurs on 2020-06-12 00:00:00

Lowest trade activity count:232 occurs on 2020-05-05 00:00:00



7 What potato price trends do you see?

Generally speaking you can use seasonality to help predict potato prices as it seems, there are times,

when the price increases drastically for the entire market (Jun/July) and it drops significantly in May.

Below you can view a plot for the weekly average potato prices across the sector.

```

In [9]: # Add a millisecond to each date to create a unique index. This allows us to turn the pr
positions_merged['Date'] = positions_merged['Date'] + \
        pd.to_timedelta(positions_merged.groupby('Date').cumcount(),
price_trends = positions.set_index('Date') # Overall market for poato trading

```

```

# Compute summary statistics
max_price = price_trends.Price.max()
min_price = price_trends.Price.min()
max_date = price_trends.Price.idxmax()
min_date = price_trends.Price.idxmin()

print("Max Price of {} occurs on {}".format(max_price, max_date))
print("Min Price of {} occurs on {}".format(min_price, min_date))

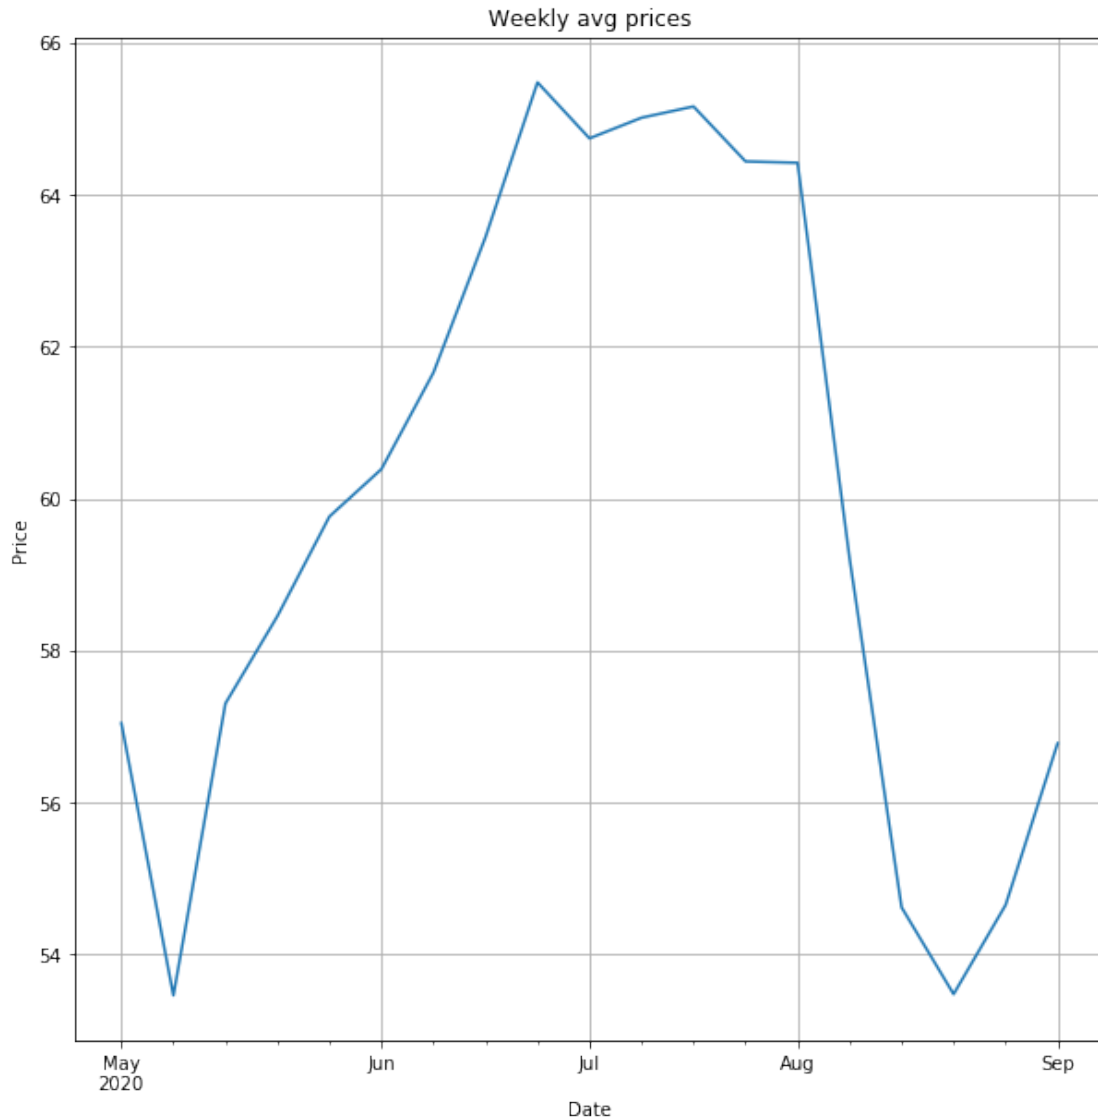
# Daily data is hard to draw a conclusion on, weekly gives a better picture of the price
weekly = price_trends['Price'].resample('W').mean()
weekly.plot(grid = True, figsize=(10,10),title='Weekly avg prices')
plt.ylabel("Price")

```

Max Price of 1519.7 occurs on 2020-08-31 00:00:00

Min Price of 8.5 occurs on 2020-08-20 00:00:00

Out[9]: Text(0, 0.5, 'Price')



8 Are there any factors that can help predict potato prices?

Yes the attributes of the potato can be used to help predict potato prices, although, not very well. The highest correlation found is the colour of the skin. Which indicates this can be used to help predict pricing using a learning algorithm.

The maturity of the potato is very good indicator of the price. Potato prices have some of the highest prices, where the maturity is shorter. Longer maturities tend to have a stable price. We can further build on this by categorizing the maturities and computing the correlation between the price. This will help us assign weights to each maturity in order to compute the correlation. Because of time constraints I will leave the exercise for future enhancements.

Below I have the correlation matrix for the attributes of the potato. I categorized the columns

using the built-in method, this method is sometimes not efficient, because the weights are arbitrary. This method provides a good estimate in understanding the data.

```
In [73]: potato_prices = positions.merge(potato_info, on='Product ID')
plt.figure(figsize = (10,10))
plt.xlabel('Price', fontsize=15)
plt.ylabel('Maturity', fontsize=15)
plt.scatter(potato_prices['Price'], potato_prices['Maturity'])
sns.despine
potato_prices.boxplot('Price', 'Maturity')

categorize_prices = potato_prices.copy()
categorize_prices['Colour of skin'] = categorize_prices['Colour of skin'].astype('category')
categorize_prices['Country'] = categorize_prices['Country'].astype('category').cat.codes
categorize_prices['Shape of tuber'] = categorize_prices['Shape of tuber'].astype('category')
categorize_prices['Colour of base of lightsprout'] = categorize_prices['Colour of base of lightsprout'].astype('category')
categorize_prices['Variety Name'] = categorize_prices['Variety Name'].astype('category')
categorize_prices['Colour of base of lightsprout'] = categorize_prices['Depth of eyes'].astype('category')
categorize_prices['Height of plants'] = categorize_prices['Height of plants'].astype('category')
categorize_prices['Frequency of berries'] = categorize_prices['Frequency of berries'].astype('category')
categorize_prices.corr()
```

```
Out [73]:
```

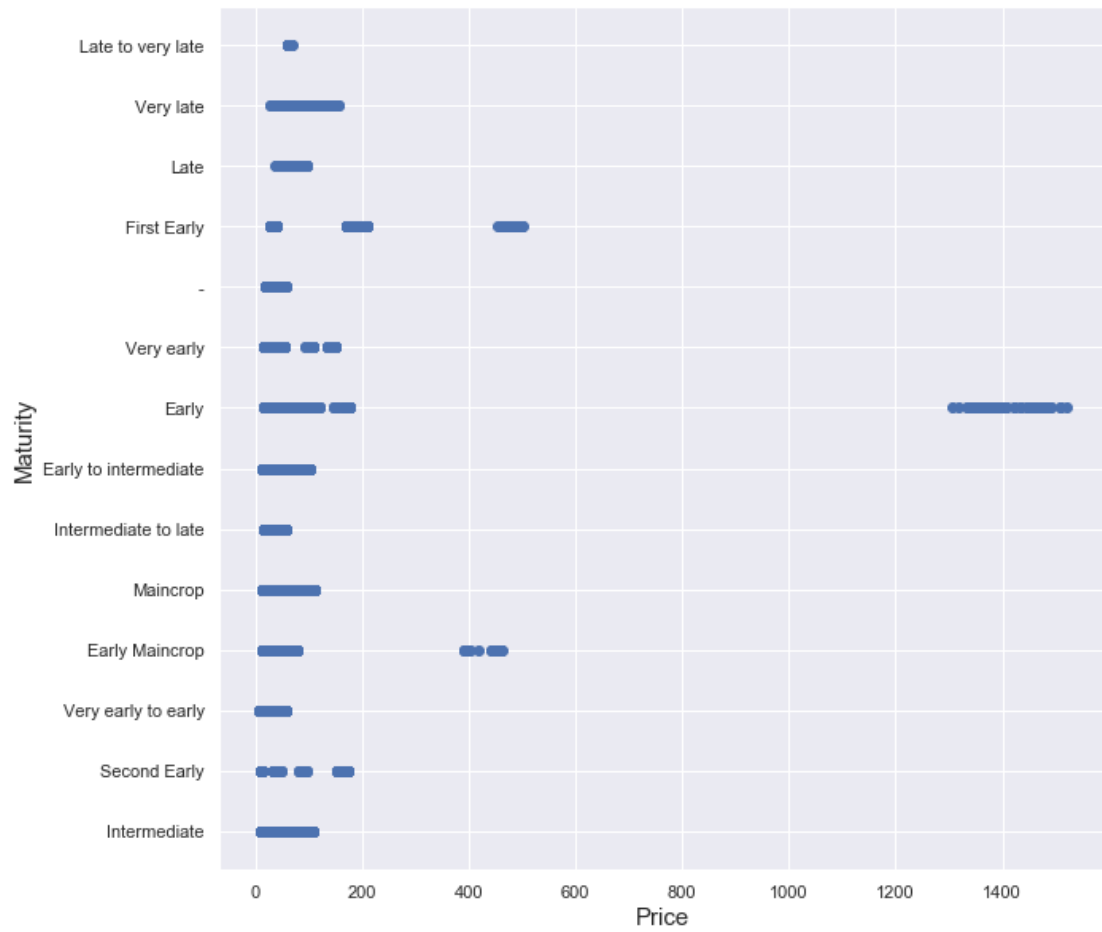
	Client ID	Product ID	Price	Variety Name	\
Client ID	1.000000	-0.074444	-0.003412	-0.075326	
Product ID	-0.074444	1.000000	0.036308	0.999431	
Price	-0.003412	0.036308	1.000000	0.035340	
Variety Name	-0.075326	0.999431	0.035340	1.000000	
Country	0.018168	0.134160	-0.042057	0.136839	
Shape of tuber	0.031648	-0.100794	-0.041551	-0.096012	
Colour of skin	-0.040825	0.109408	0.164066	0.103298	
Colour of base of lightsprout	-0.027587	0.108555	0.091351	0.105588	
Height of plants	0.034261	0.267695	0.031578	0.273255	
Frequency of berries	-0.000672	-0.162083	-0.121572	-0.155670	

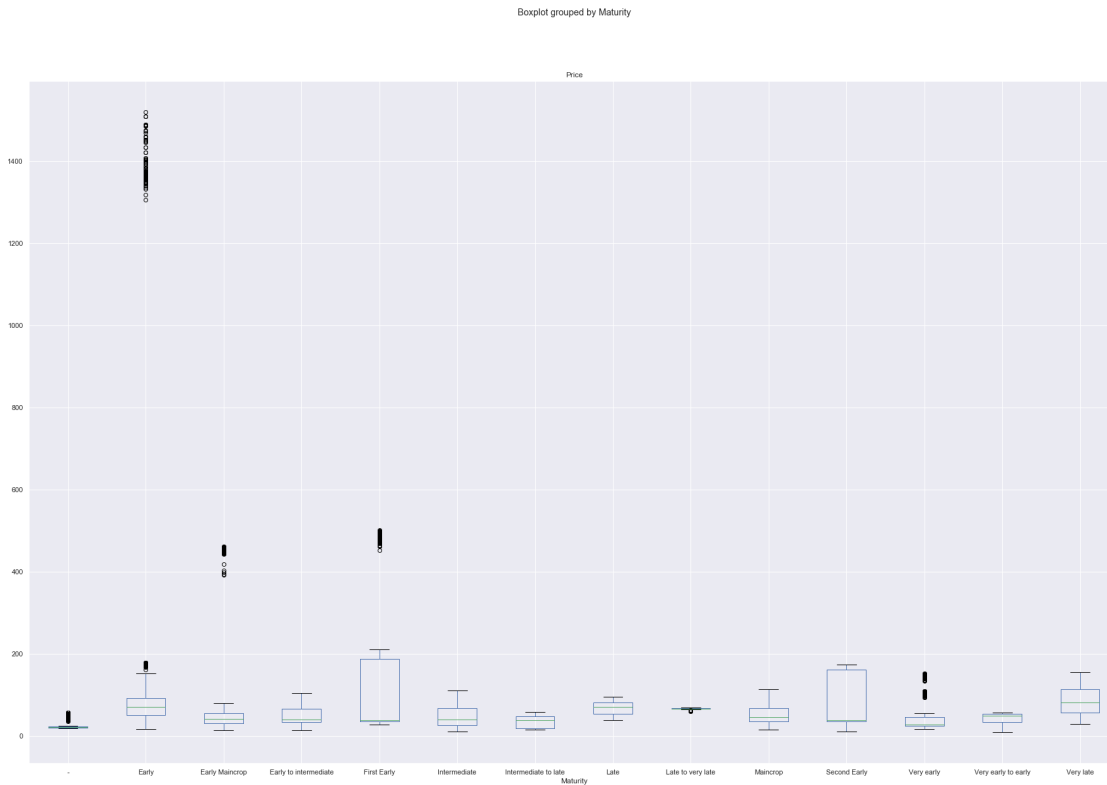
	Country	Shape of tuber	Colour of skin	\
Client ID	0.018168	0.031648	-0.040825	
Product ID	0.134160	-0.100794	0.109408	
Price	-0.042057	-0.041551	0.164066	
Variety Name	0.136839	-0.096012	0.103298	
Country	1.000000	-0.065473	0.052011	
Shape of tuber	-0.065473	1.000000	-0.040371	
Colour of skin	0.052011	-0.040371	1.000000	
Colour of base of lightsprout	0.058634	0.039825	0.138018	
Height of plants	0.235765	-0.054195	0.107843	
Frequency of berries	0.217766	0.045332	-0.218283	

	Colour of base of lightsprout	\
Client ID	-0.027587	

Product ID	0.108555
Price	0.091351
Variety Name	0.105588
Country	0.058634
Shape of tuber	0.039825
Colour of skin	0.138018
Colour of base of lightsprout	1.000000
Height of plants	-0.022613
Frequency of berries	-0.024054

	Height of plants	Frequency of berries
Client ID	0.034261	-0.000672
Product ID	0.267695	-0.162083
Price	0.031578	-0.121572
Variety Name	0.273255	-0.155670
Country	0.235765	0.217766
Shape of tuber	-0.054195	0.045332
Colour of skin	0.107843	-0.218283
Colour of base of lightsprout	-0.022613	-0.024054
Height of plants	1.000000	0.090708
Frequency of berries	0.090708	1.000000





9 Are there any factors that can help predict client activity?

A good indicator for client activity is their region, with a dominant portion of traders coming from clients in the midwest.

Another indicator for trading activity is the the industry where there is strong correlation between Industrial clients and trading activity.

Attributes from the potatoes can also predict the client activity. The most traded potato id can give a sense for future activity.

```
In [77]: print(positions_merged.groupby(['Client Location',])['Date'].count().sort_values(ascending=
print(positions_merged.groupby(['Client Type',])['Date'].count().sort_values(ascending=
positions_merged.groupby(['Maturity'])['Date'].count().sort_values().head()
positions_merged.groupby(['Product ID'])['Date'].count().sort_values().head()
```

Client Location

Midwest 13142

Northeast 6743

North 4038

West 1956

Name: Date, dtype: int64

Client Type

Industrial 19652

```
Organic      2316
Hydroponic   2238
Traditional  1673
Name: Date, dtype: int64
```

```
Out[77]: Product ID
        20      9
        284    12
         6     13
        280    23
        130    25
Name: Date, dtype: int64
```

10 Summarize what has happened during this period

During this period the activity and prices are their lowest in May and peak in June.

Summer months appear to be where the highest activity and prices occur for this product type.

The fact that activity and pricing is consistent is intuitive. One would reason, that as market interest increases for a particular product, the price will be driven higher.

Potato prices movements are generally in line with each other, following similar trends and have a high correlation between product ids.

Client market control is heavily favored by clients coming from the Industrial sector in the Midwest region. We can use this to help predict pricing as we know where the activity and potato pricing lies.

10.0.1 We see majority of the market value is controlled by those located in the Midwest.

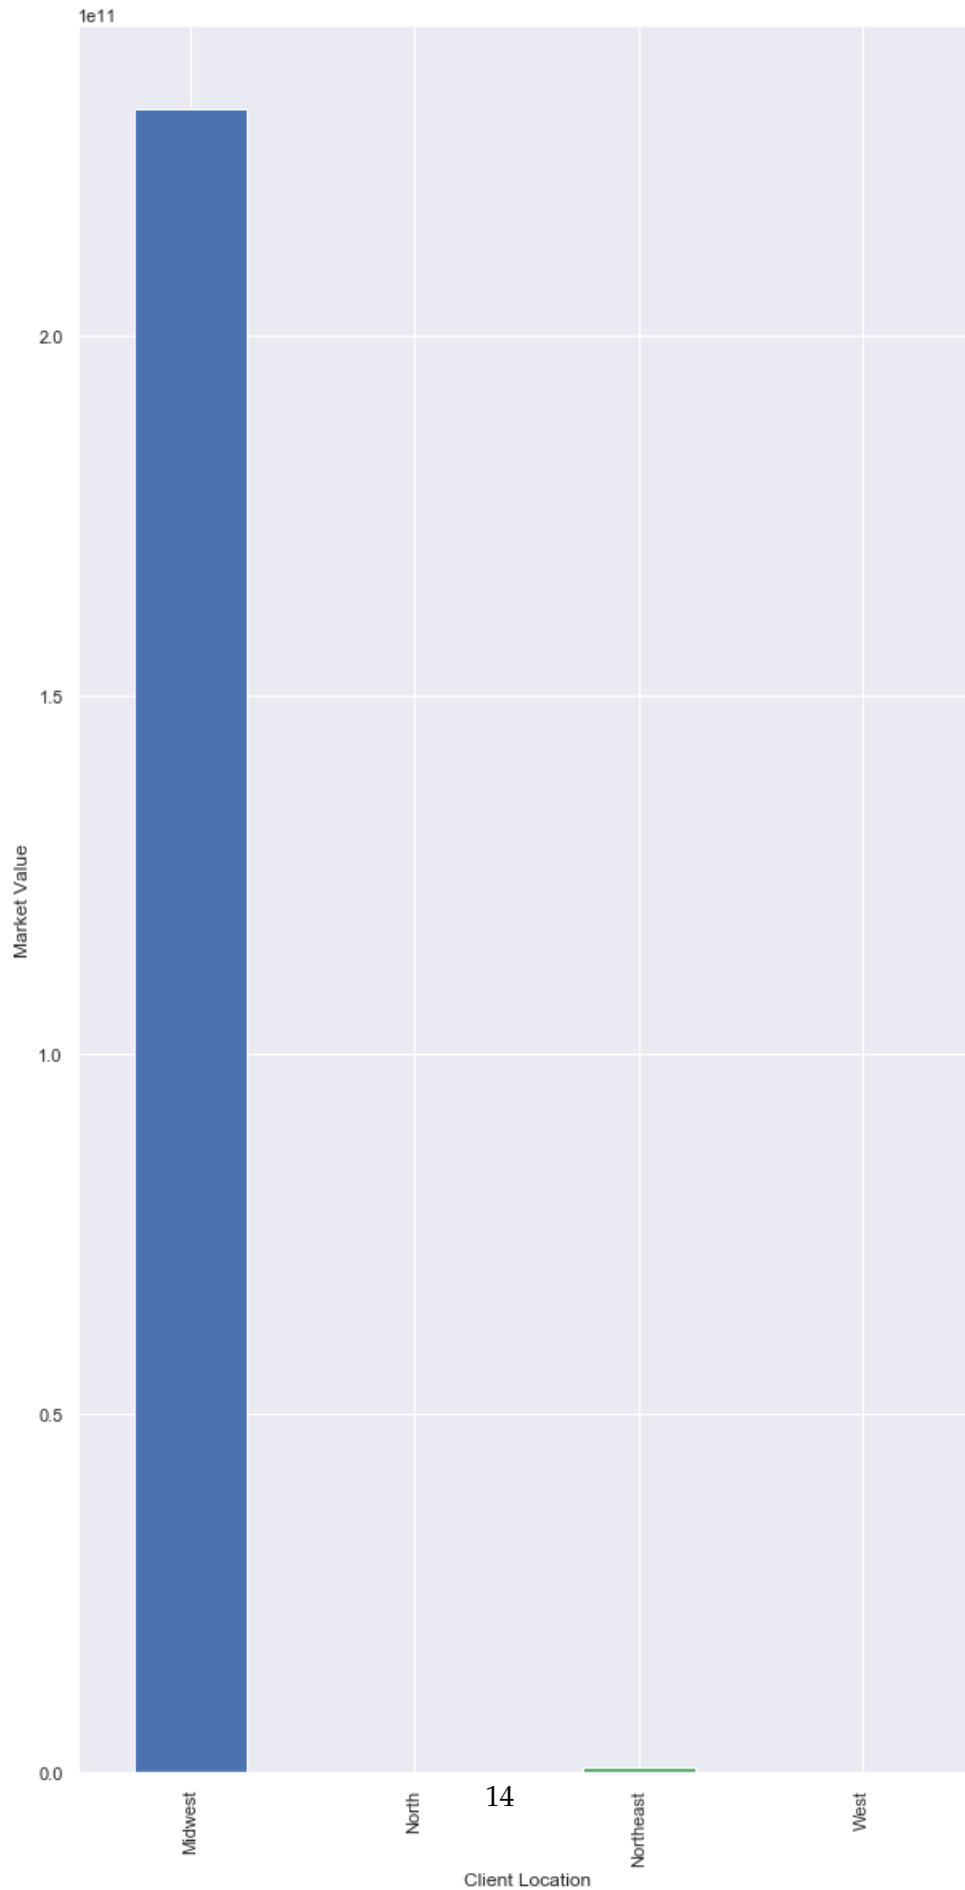
10.0.2 The next largest is the Northeast but is pale in comparison,

```
In [42]: region_sum = positions_merged.groupby('Client Location')['Market Value'].sum()
        print(region_sum.apply(lambda x: '{0:.2f}'.format(x)))
```

```
region_sum_plt = region_sum.plot(kind='bar',figsize=(10,20))
region_sum_plt.set_ylabel('Market Value')
```

```
Client Location
Midwest      231650623218.59
North        127527663.01
Northeast    764053837.21
West         95691518.80
Name: Market Value, dtype: object
```

```
Out[42]: Text(0, 0.5, 'Market Value')
```

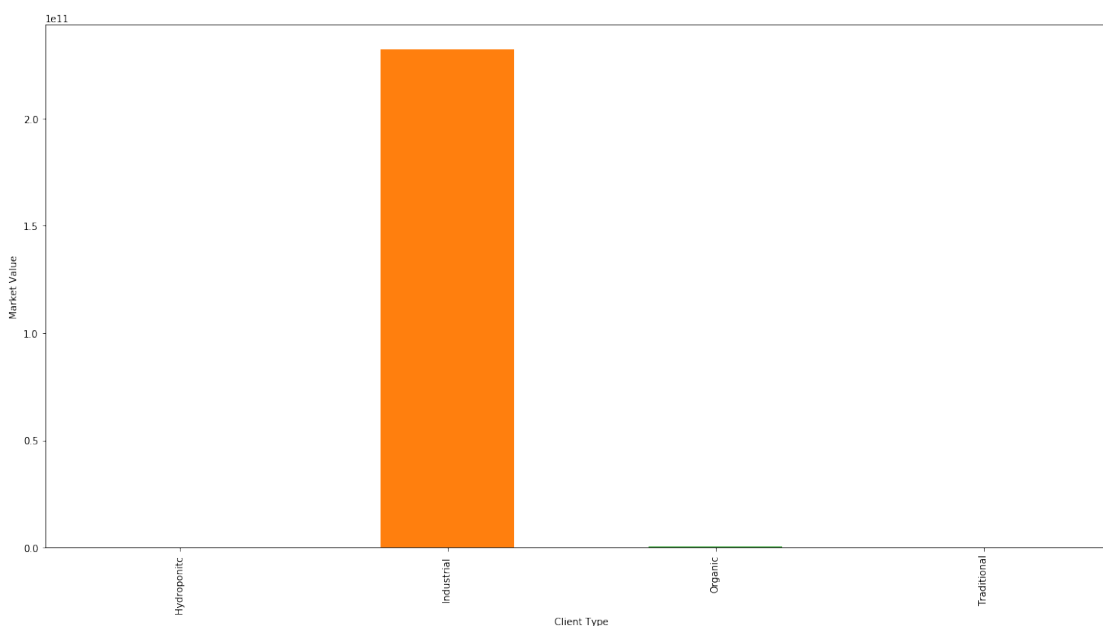


10.0.3 Majortiy of the market share is controlled by clients in the Industrial field

```
In [8]: industry_sum = positions_merged.groupby('Client Type')['Market Value'].sum()
        print(industry_sum.apply(lambda x: '{0:.2f}'.format(x)))
        industry_sum.plot(kind='bar', figsize=(20,10))
        plt.ylabel('Market Value')
```

```
Client Type
Hydroponic      54930213.16
Industrial    232056014860.38
Organic         441874498.27
Traditional      85076665.81
Name: Market Value, dtype: object
```

```
Out[8]: Text(0, 0.5, 'Market Value')
```



Pricing between products are highly correlated and follow similar trend throughout this time period.

```
In [44]: pivoted = pd.pivot_table(price_trends, values='Price', columns='Product ID', index='Date')
        pivoted.idxmax(1).to_frame().T

        pivoted.plot(grid = True, figsize=(20,20))
        plt.ylabel("Price")
        #print(pivoted.corr())
```

```
Out[44]: Text(0, 0.5, 'Price')
```